INTRODUCTION TO QUERY STORE

SQLPORT 116, ONLINE MEETING

MILOŠ RADIVOJEVIĆ, PRINCIPAL DATABASE CONSULTANT, BWIN GVC, AUSTRIA

ABOUT ME

- MILOŠ RADIVOJEVIĆ
- Data Platform MVP
- Principal Database Consultant at bwin, Vienna, Austria
- Co-Founder: SQL Pass Austria
- Conference Speaker, Book Author

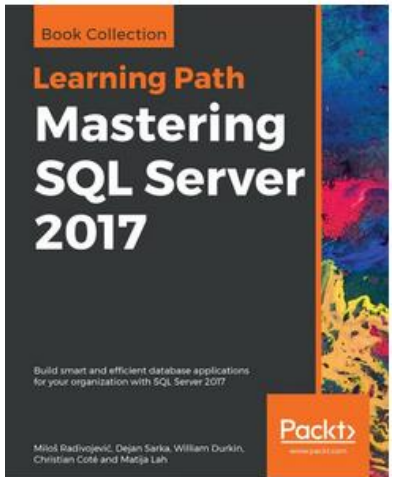

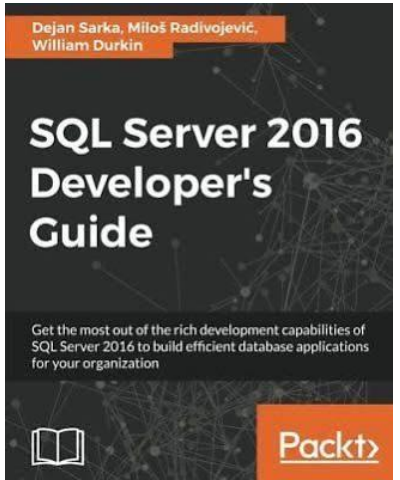

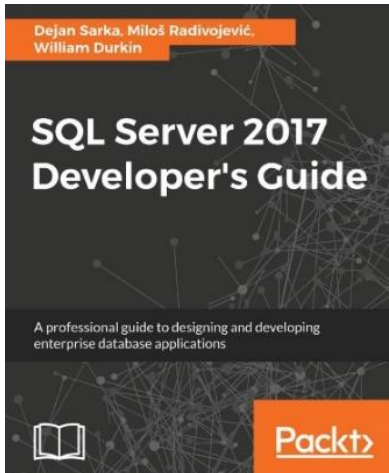

AGENDA

- Life before Query Store
- Why Query Store?
- Query Store Architecture and Configuration
- Query Store in Action

PROBLEMS WITH DMV TROUBLESHOOTING

- reflects particular or aggregated information from the last server restart only
- when server goes down, cache disappears
- only the latest execution plan for a query is available
- you don’t know if and how execution plan changed over time

MY QUERY STORY

MY QUERY STORY

- Special stored procedure
  - Collect exec. parameters for each stored procedure (total elapsed time, total logical reads) within a database
  - Collect all different execution plans (XML) for the most important stored procedures
  - Store collected results in tables within a database
- SQL Job
- Call the special SP every 5 minutes

MY QUERY STORY

```
SELECT TOP 100 *, DATALENGTH(fQueryPlan) AS L FROM [illegible] WHERE frId = 188 ORDER BY 1 DESC
```

100 %

Results | Messages

	frId	fSnapshotDate	fExecutionCount	fTotalElapsedTime	fTotalLogicalReads	fCachedTime	fQueryPlan	L
1	188	2015-10-06 13:00:01.270	75451	3010807126	1556870293	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
2	188	2015-10-06 12:00:01.363	74971	2991984410	1547141419	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
3	188	2015-10-06 11:00:02.253	74491	2973125700	1537429268	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
4	188	2015-10-06 10:00:01.363	74011	2954006611	1527718630	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
5	188	2015-10-06 09:00:01.390	73531	2935445234	1518024064	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
6	188	2015-10-06 08:00:02.143	73051	2915978882	1508318816	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26101
7	188	2015-10-06 07:00:01.503	72571	2894914949	1498608189	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
8	188	2015-10-06 06:00:01.597	72091	2875729099	1488895815	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
9	188	2015-10-06 05:00:01.323	71611	2857248974	1479176199	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
10	188	2015-10-06 04:00:01.340	71131	2839315557	1469447430	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
11	188	2015-10-06 03:00:01.370	70651	2819743677	1459718314	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
12	188	2015-10-06 02:00:01.827	70171	2797400153	1449987319	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
13	188	2015-10-06 01:00:01.983	69692	2775447825	1440277391	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
14	188	2015-10-06 00:00:01.547	69213	2756110848	1430569882	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095
15	188	2015-10-05 23:00:01.603	68733	2737765742	1420828904	2015-09-30 03:05:11.590	<ShowPlanXML xmlns="http://schemas.microsoft.com...	26095

MY QUERY STORY

Row Labels	Sum of AvgElapsedMSPreviousWeek	Sum of AvgElapsedMS
30.09.2015	49	78
30.09.2015 01:00	66	92
30.09.2015 02:00	92	107
30.09.2015 04:00	71	51
30.09.2015 05:00	65	50
30.09.2015 06:00		
30.09.2015 07:00		
30.09.2015 08:00		
30.09.2015 09:00		
30.09.2015 10:00		
30.09.2015 11:00		
30.09.2015 12:00		
30.09.2015 13:00		
30.09.2015 14:00		
30.09.2015 15:00		
30.09.2015 16:00		
30.09.2015 17:00		
30.09.2015 18:00		
30.09.2015 19:00		
30.09.2015 20:00		
30.09.2015 21:00		
30.09.2015 22:00		
30.09.2015 23:00		
Grand Total		

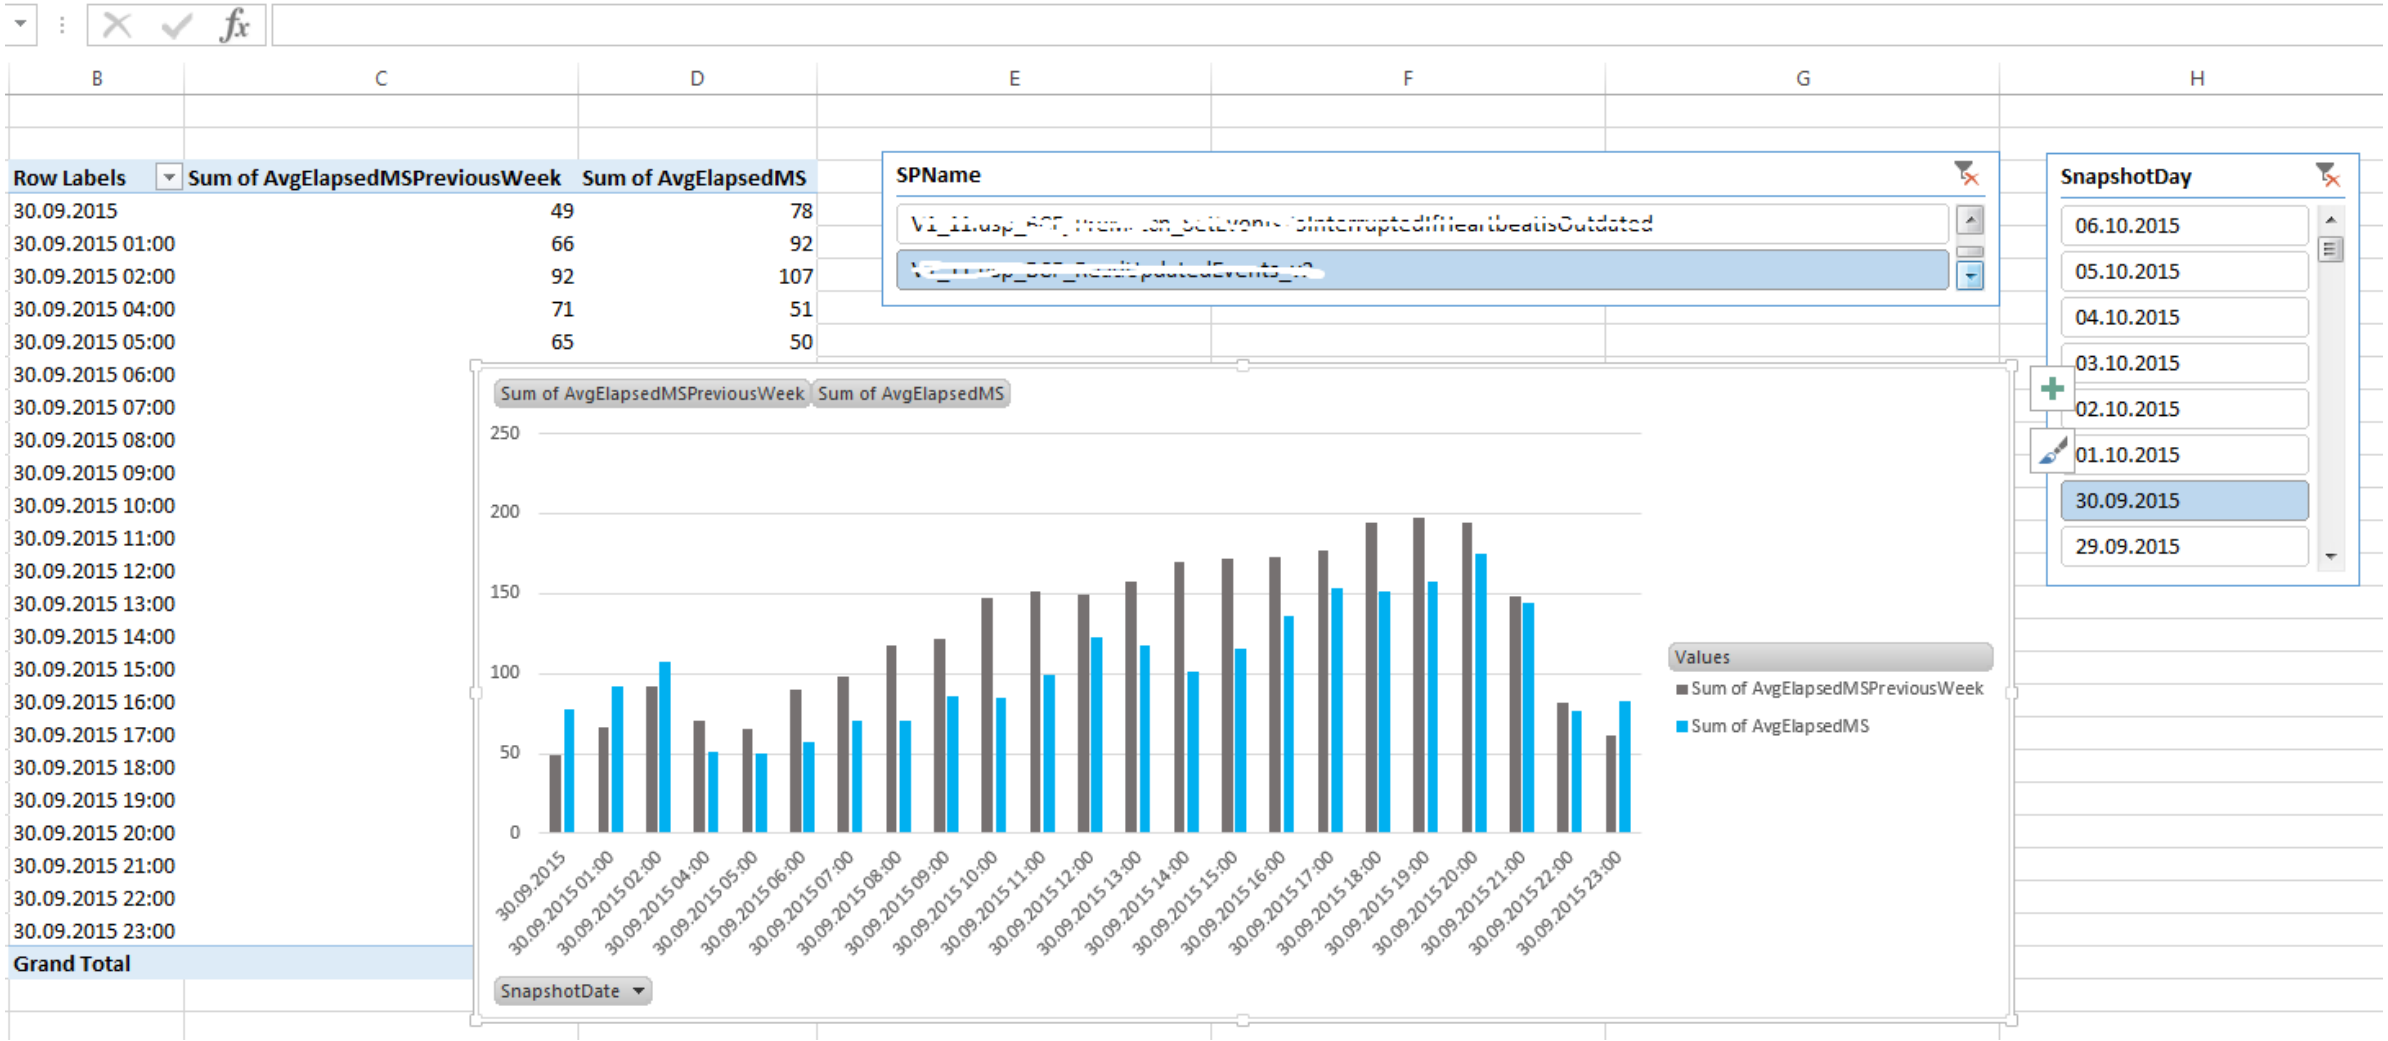

MY QUERY STORY

Row Labels	Sum of AvgElapsedMSPreviousWeek	Sum of AvgElapsedMS
05.09.2015	14	50
05.09.2015 01:00	19	67
05.09.2015 02:00	24	104
05.09.2015 03:00	30	85
05.09.2015 04:00	31	52
05.09.2015 06:00	12	46
05.09.2015 07:00	15	54
05.09.2015 08:00	18	
05.09.2015 09:00	18	
05.09.2015 10:00	20	
05.09.2015 11:00	18	
05.09.2015 12:00	29	
05.09.2015 13:00	25	
05.09.2015 14:00	25	
05.09.2015 15:00	34	
05.09.2015 16:00	33	
05.09.2015 17:00	42	
05.09.2015 18:00	46	
05.09.2015 19:00	46	
05.09.2015 20:00	38	
05.09.2015 21:00	36	
05.09.2015 22:00	35	
05.09.2015 23:00	30	
Grand Total	638	

SPName

SnapshotDay

10.09.2015
09.09.2015
08.09.2015
07.09.2015
06.09.2015
05.09.2015
04.09.2015
03.09.2015

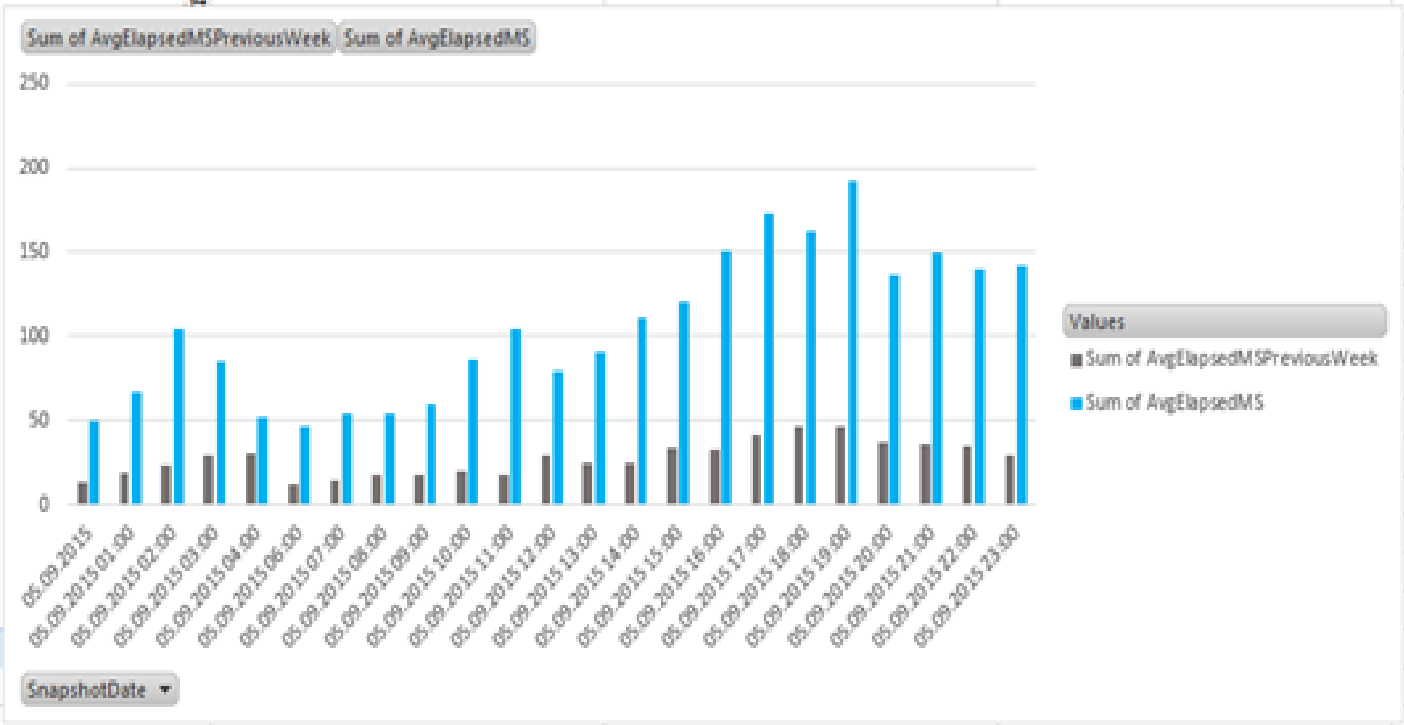

WHAT DOES THE QUERY STORE?

- Stores the history of queries, plans, execution details and waits statistics
  - Stores also plans with OPTION (RECOMPILE)
- Belongs to database
  - it is persistent – survives after restart, failover etc.
- Introduced with SQL Server 2016
- Disabled by default in on-prem databases, in SQL Azure is enabled
- Available in all editions (a few Enterprise features)
- Requires VIEW_DATABASE_STATE permission
- Cannot be enabled for system databases

QUERY STORE ARCHITECTURE

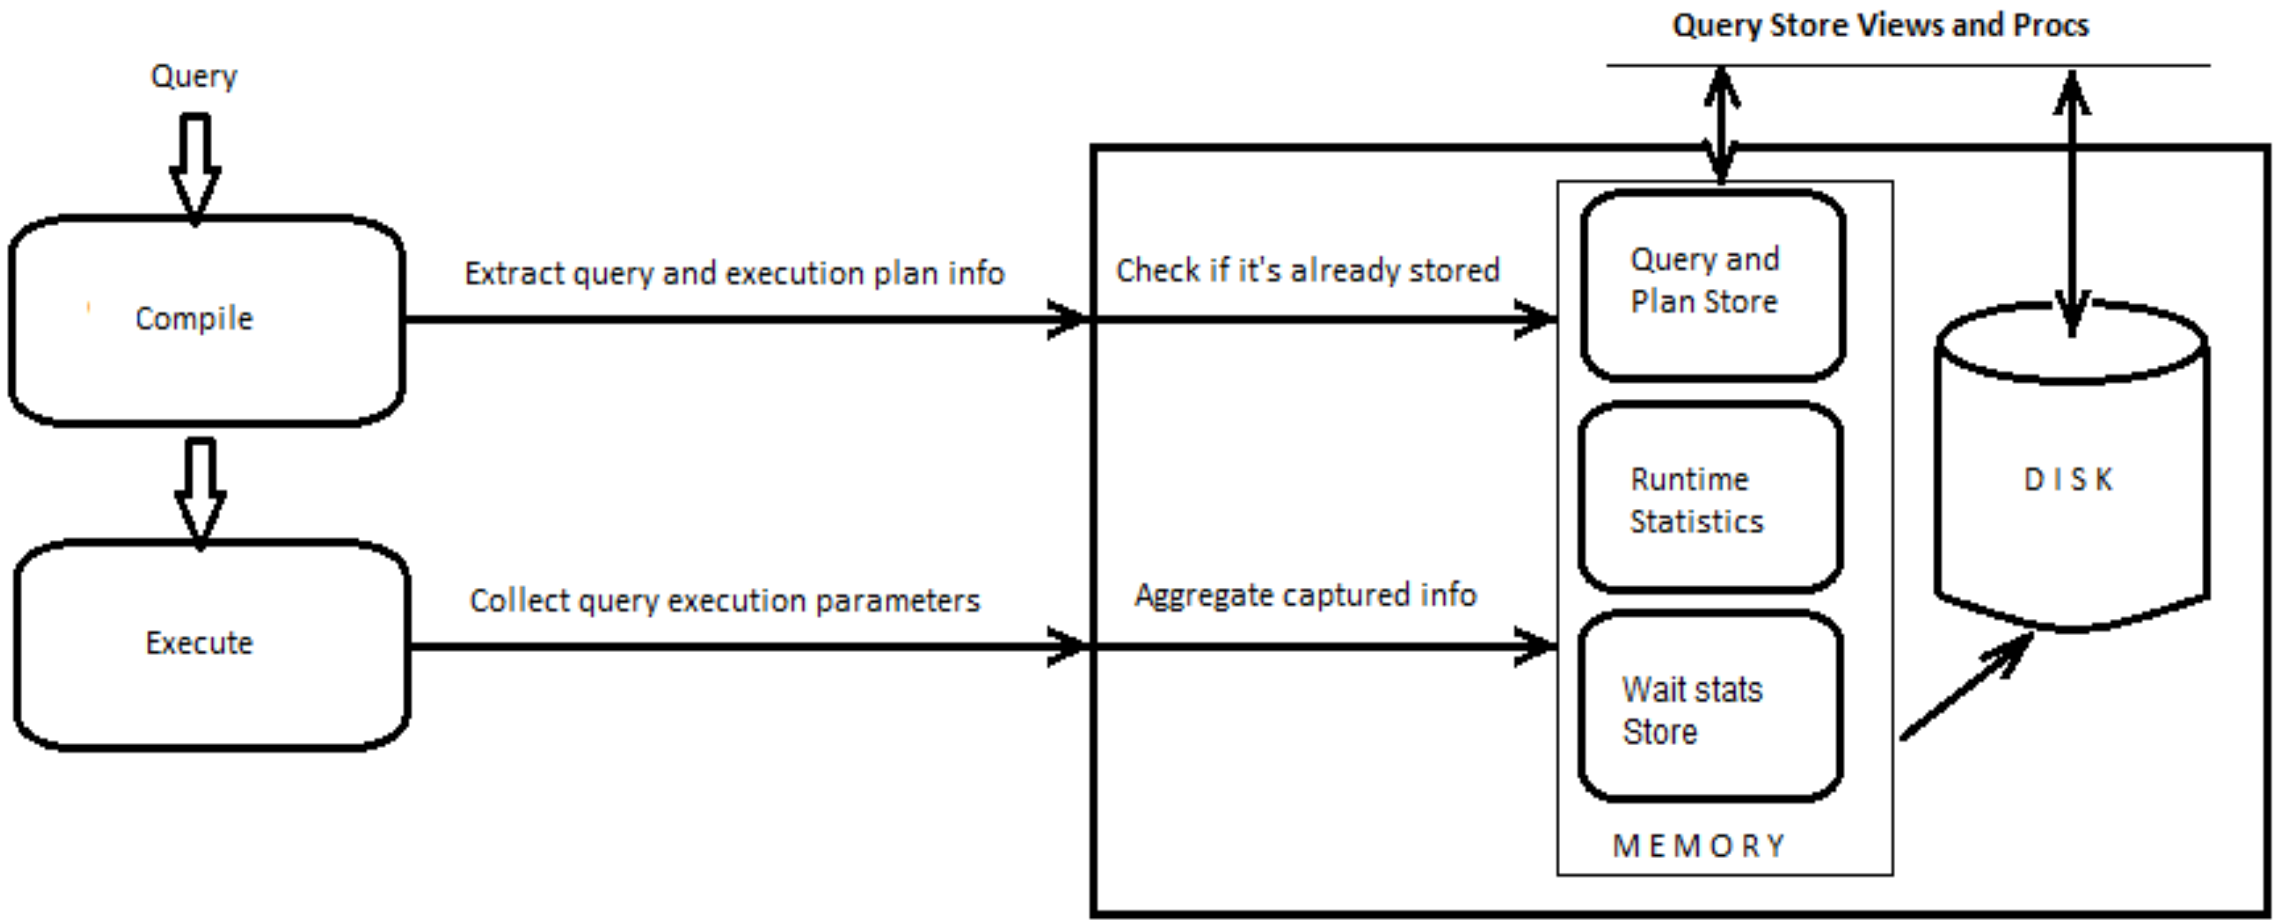

QUERY STORE

Query Text

Query
(id, object_id, compile details)

Execution Plans

Execution Parameters
(duration, CPU time, logical reads...)

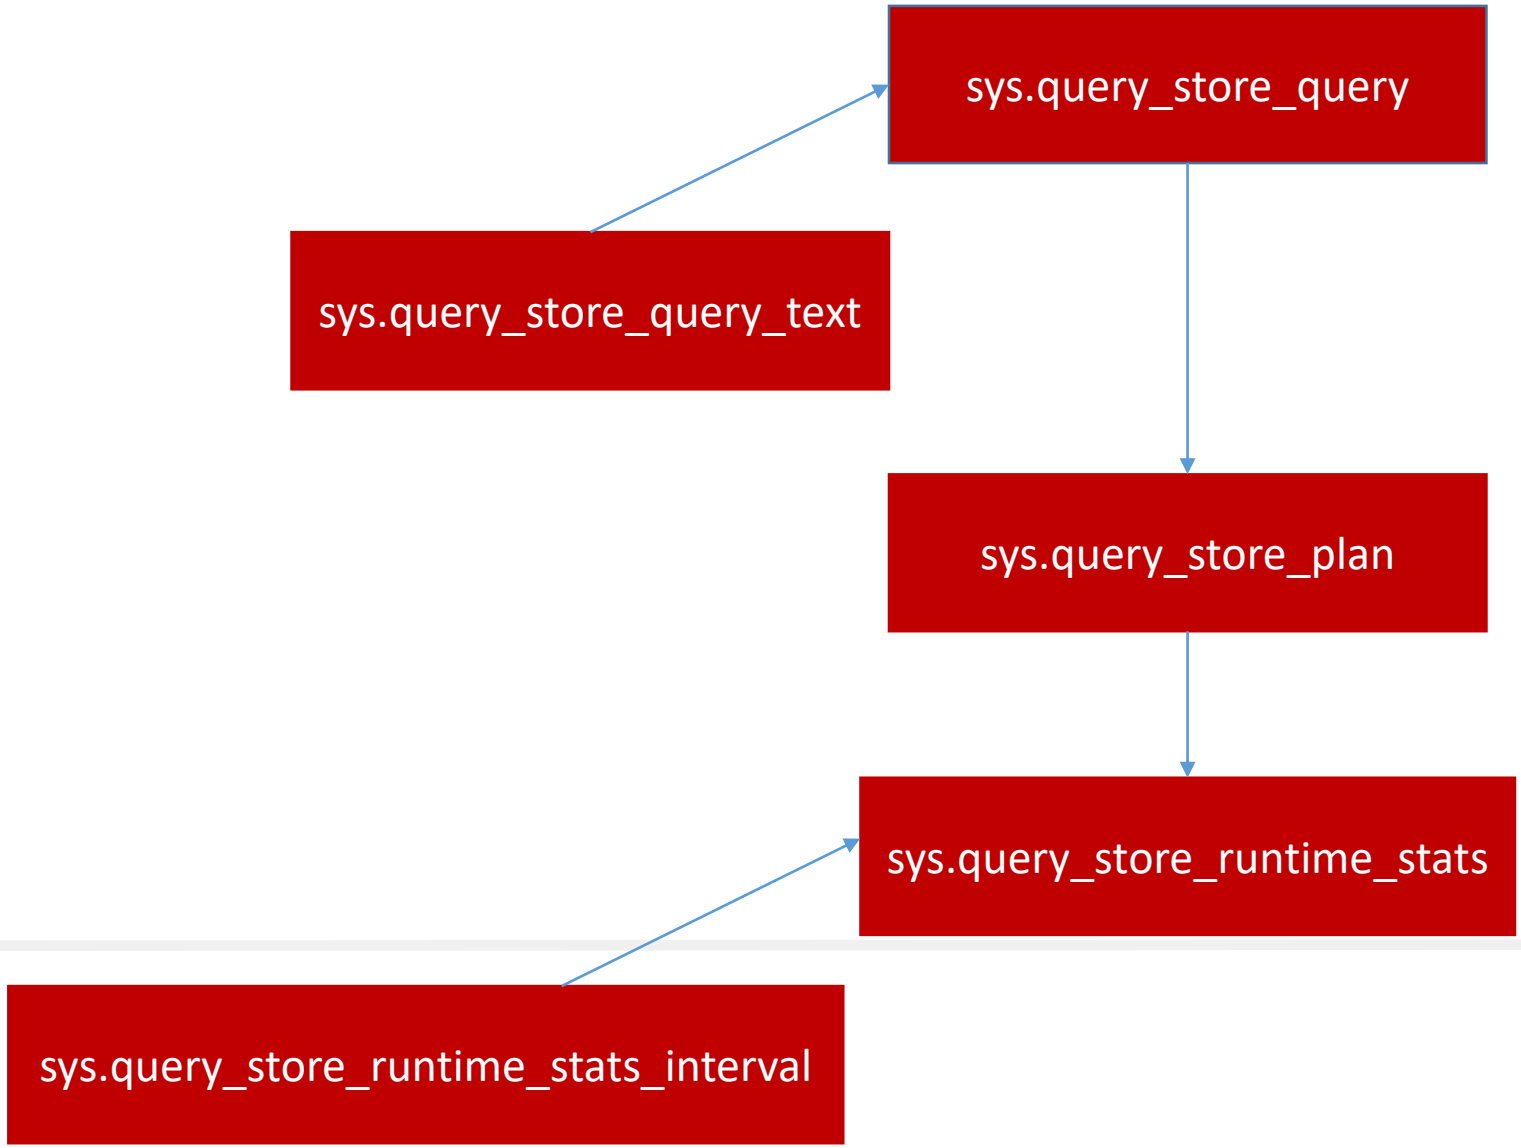

WHAT CAN WE DO WITH QUERY STORE?

- Identifying query plan regressions
- Fixing plan regression (by forcing plans)
- Reducing the risk of upgrading, patching and reconfiguring
- Analyzing workload patterns
- Understanding waits at the plan level
- Identifying the most expensive and queries degraded over time

WHAT CAN WE DO WITH QUERY STORE?

- We can have answers to the following questions:
  - Was this query slow last weekend?
  - Why my query was slow last Saturday?
  - What are unstable queries (with multiple plans)?
  - We have some timeouts in the application. Are they from database?
  - Find out unfinished queries or queries that ended with an exception

CONFIGURING QUERY STORE

```
ALTER DATABASE WideWorldImporters SET QUERY_STORE = ON;
```

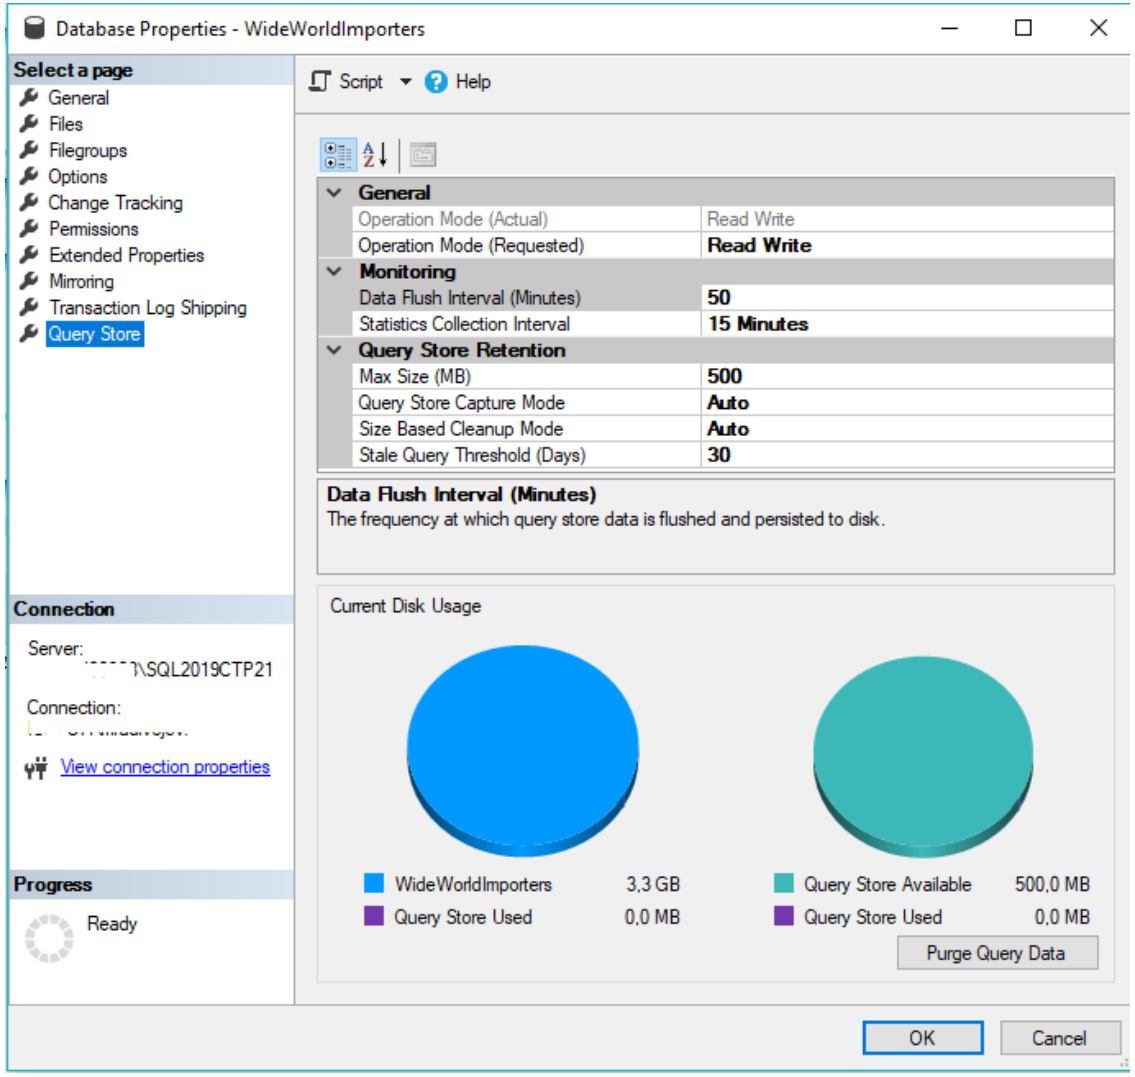

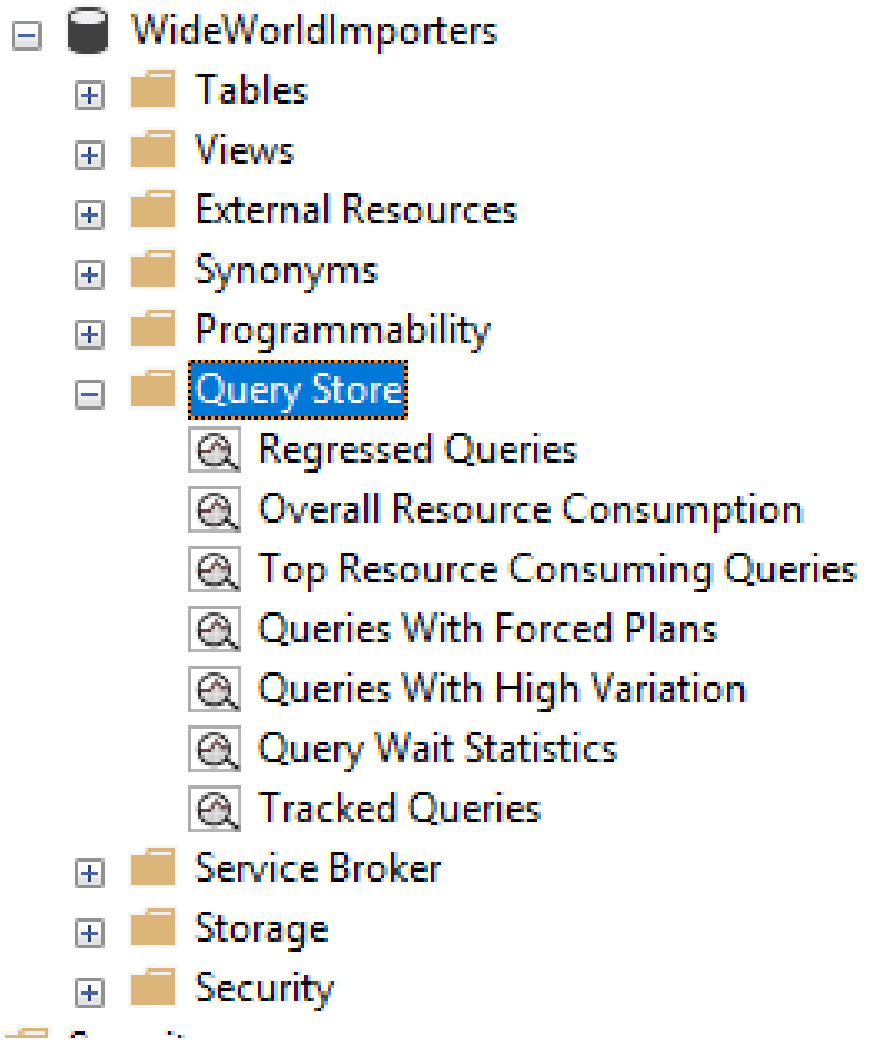

CONFIGURING QUERY STORE

```
ALTER DATABASE WideWorldImporters SET QUERY_STORE = ON (
OPERATION_MODE = READ_WRITE,
CLEANUP_POLICY = (STALE_QUERY_THRESHOLD_DAYS = 30),
DATA_FLUSH_INTERVAL_SECONDS = 900,
INTERVAL_LENGTH_MINUTES = 60,
MAX_STORAGE_SIZE_MB = 1000,
QUERY_CAPTURE_MODE = AUTO,
SIZE_BASED_CLEANUP_MODE = AUTO,
MAX_PLANS_PER_QUERY = 200,
WAIT_STATS_CAPTURE_MODE = ON,
)
```

CONFIGURING QUERY STORE

- OPERATION_MODE = { READ_WRITE | READ_ONLY | OFF | ERROR }
- CLEANUP_POLICY = (STALE_QUERY_THRESHOLD_DAYS = 30),
- DATA_FLUSH_INTERVAL_SECONDS = 900,
- INTERVAL_LENGTH_MINUTES = {1 | 5 | 10 | 15 | 30 | 60 | 1440}
- MAX_STORAGE_SIZE_MB = 1000,
- QUERY_CAPTURE_MODE = { ALL | AUTO | CUSTOM | NONE}
- SIZE_BASED_CLEANUP_MODE = { AUTO | OFF }
- MAX_PLANS_PER_QUERY = 200,
- WAIT_STATS_CAPTURE_MODE = { ON | OFF }

CONFIGURING QUERY STORE

```
ALTER DATABASE WideWorldImporters SET QUERY_STORE = ON (
QUERY_CAPTURE_MODE = CUSTOM,
      QUERY_CAPTURE_POLICY = (
      EXECUTION_COUNT = 30,
      TOTAL_COMPILE_CPU_TIME_MS = 1000,
      TOTAL_EXECUTION_CPU_TIME_MS = 100,
      STALE_CAPTURE_POLICY_THRESHOLD = 24 HOURS
    )
)
```

QUERY STORE IN ACTION

PERFORMANCE IMPACT

- 3-5%
- 4 new Query Store perf counters
- Extended Events:
  - SQL 2017 – 91 event
  - SQL 2019 – 92 events

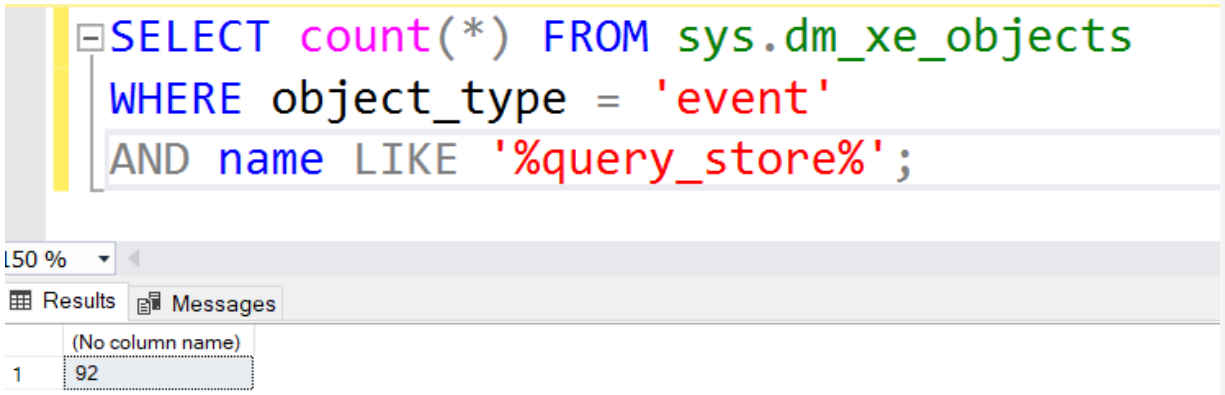

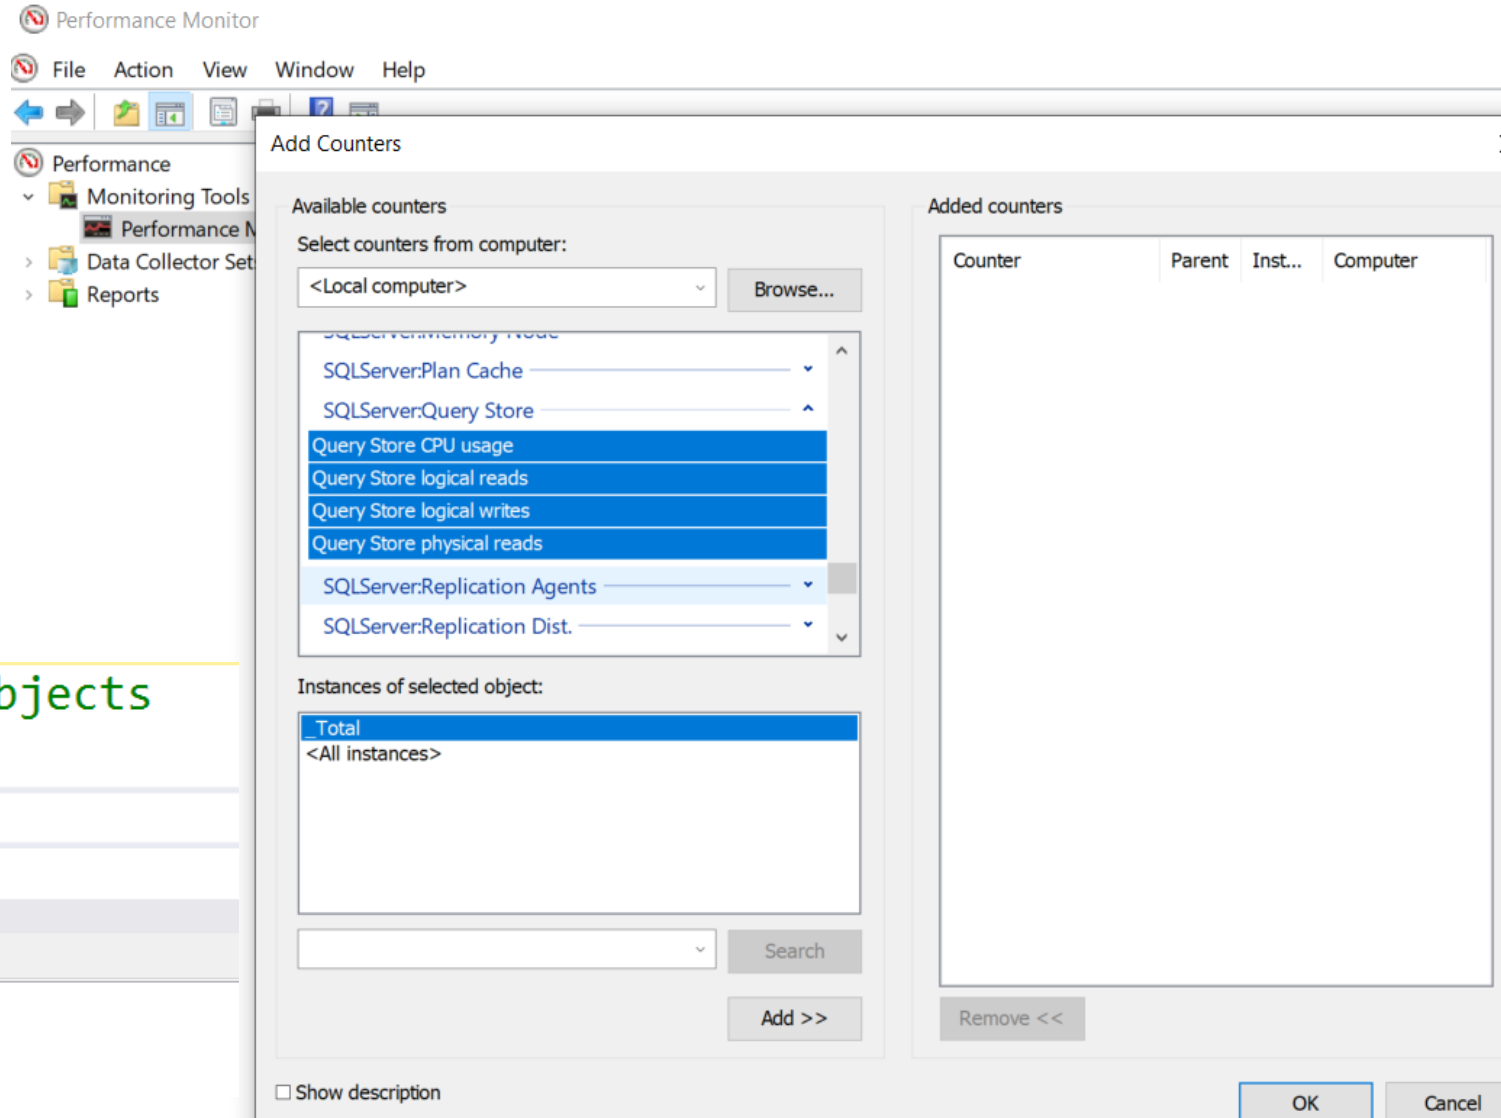

PLAN FORCING

- Should be considered as mitigation and not a solution
  - Buying time to fix the query
  - Not always a good idea!
  - Check carefully plan history for a *query_id* before you force the plan
- Be careful with parameter sensitive queries
  - Plan forcing is NOT a solution for parameter sniffing!
- Use **sp_query_store_force_plan** rather than GUI
- Requires **db_owner** permissions
- You cannot force a plan for query if it hasn't been generated for that query
- You cannot force the plan if Query Store is in the *OFF* operation mode (in *Read-Only* is possible)

WHICH PLAN TO FORCE?

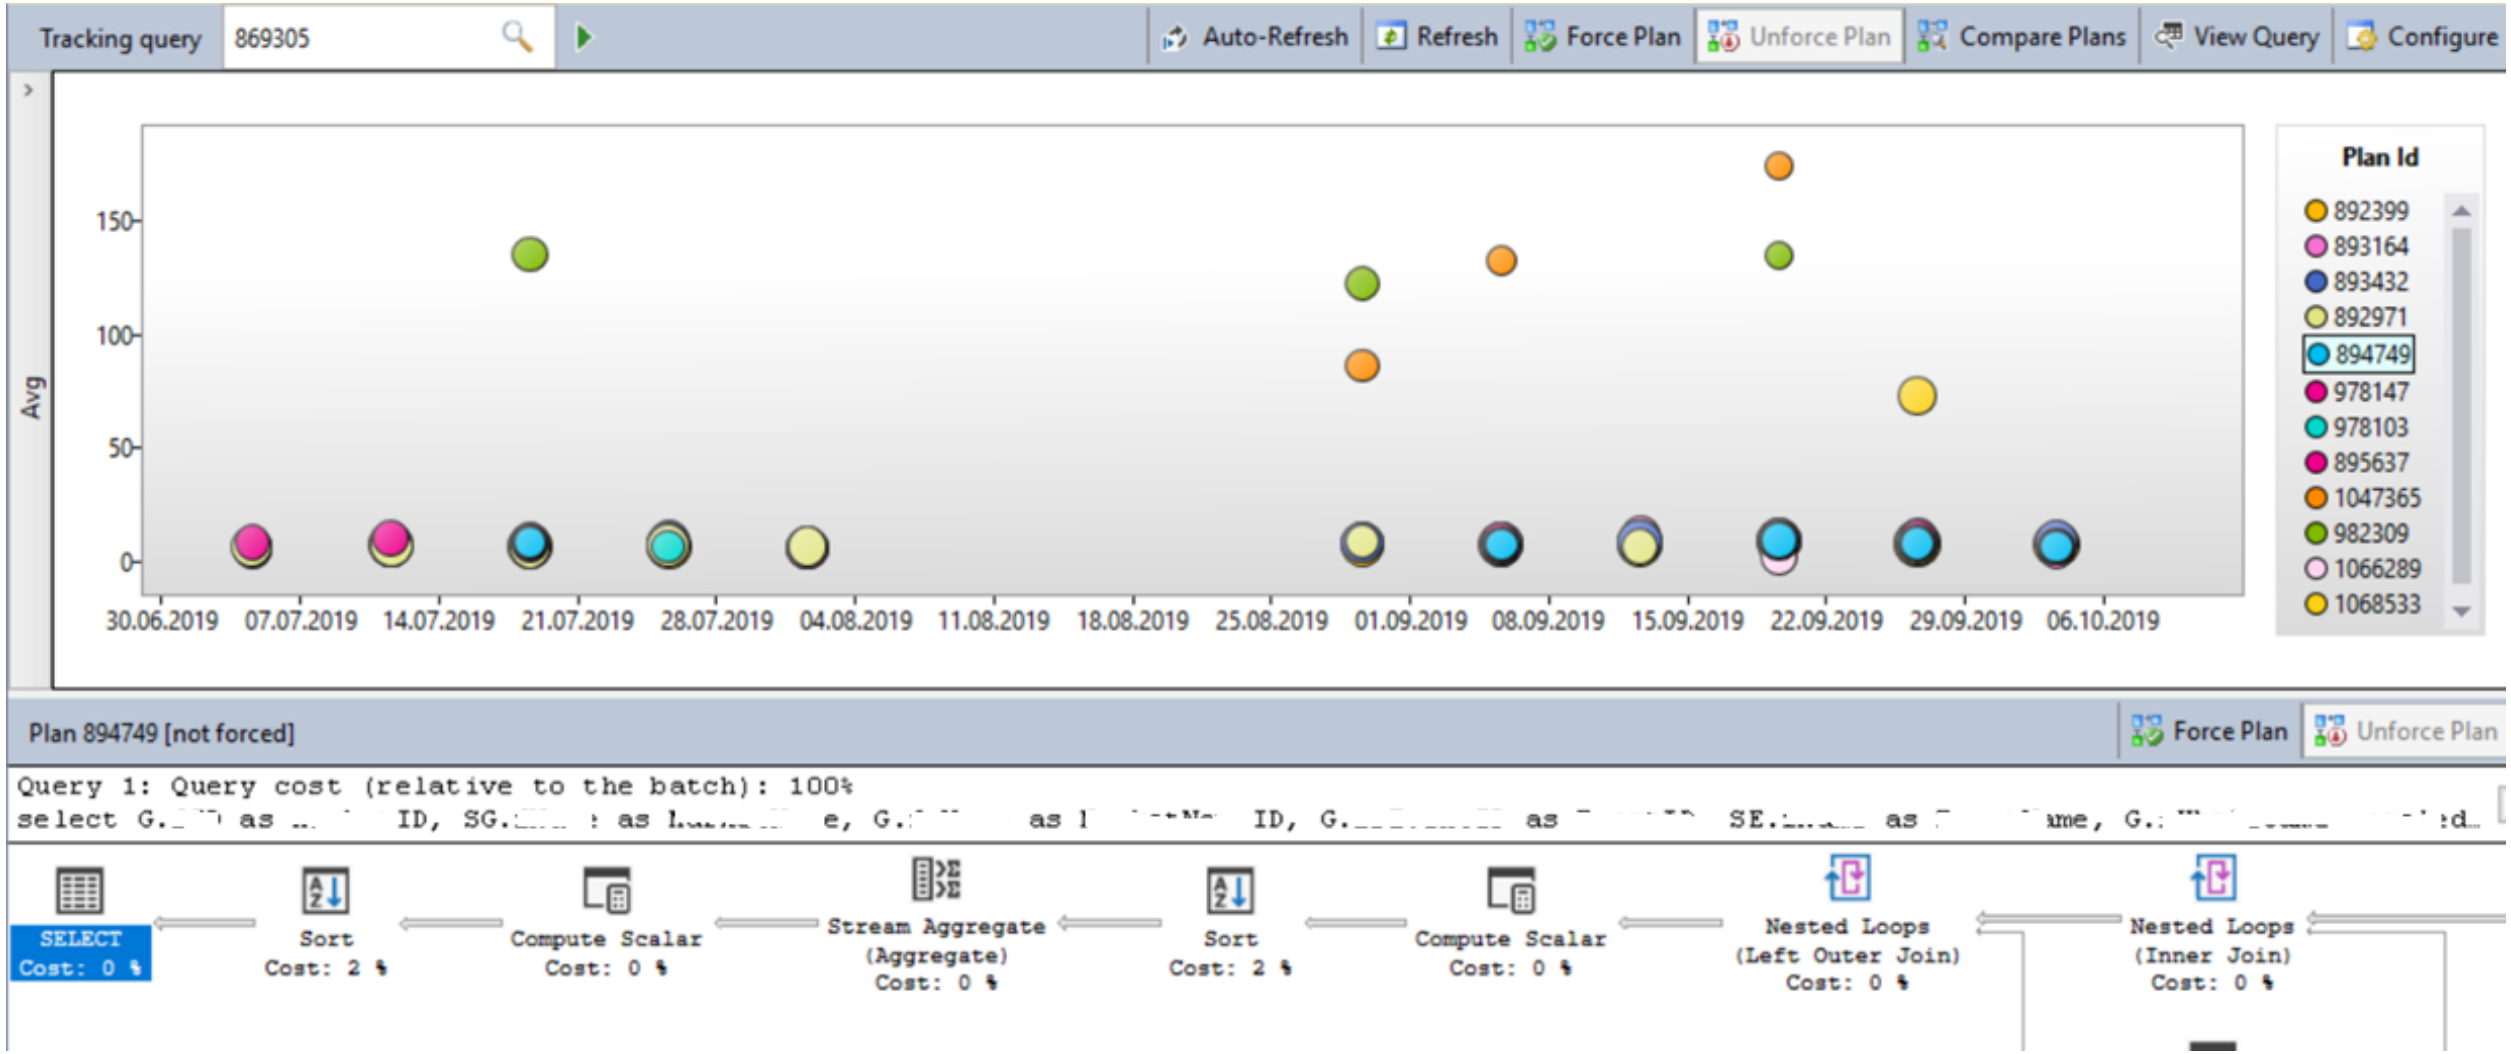

QUERY STORE REPORTS

- Regressed Queries
- Overall Resource Consumption
- Top Resource Consuming Queries
- Queries With Forced Plans
- Queries With High Variation
- Query Waits Statistics
- Tracked Queries

REGRESSED QUERIES

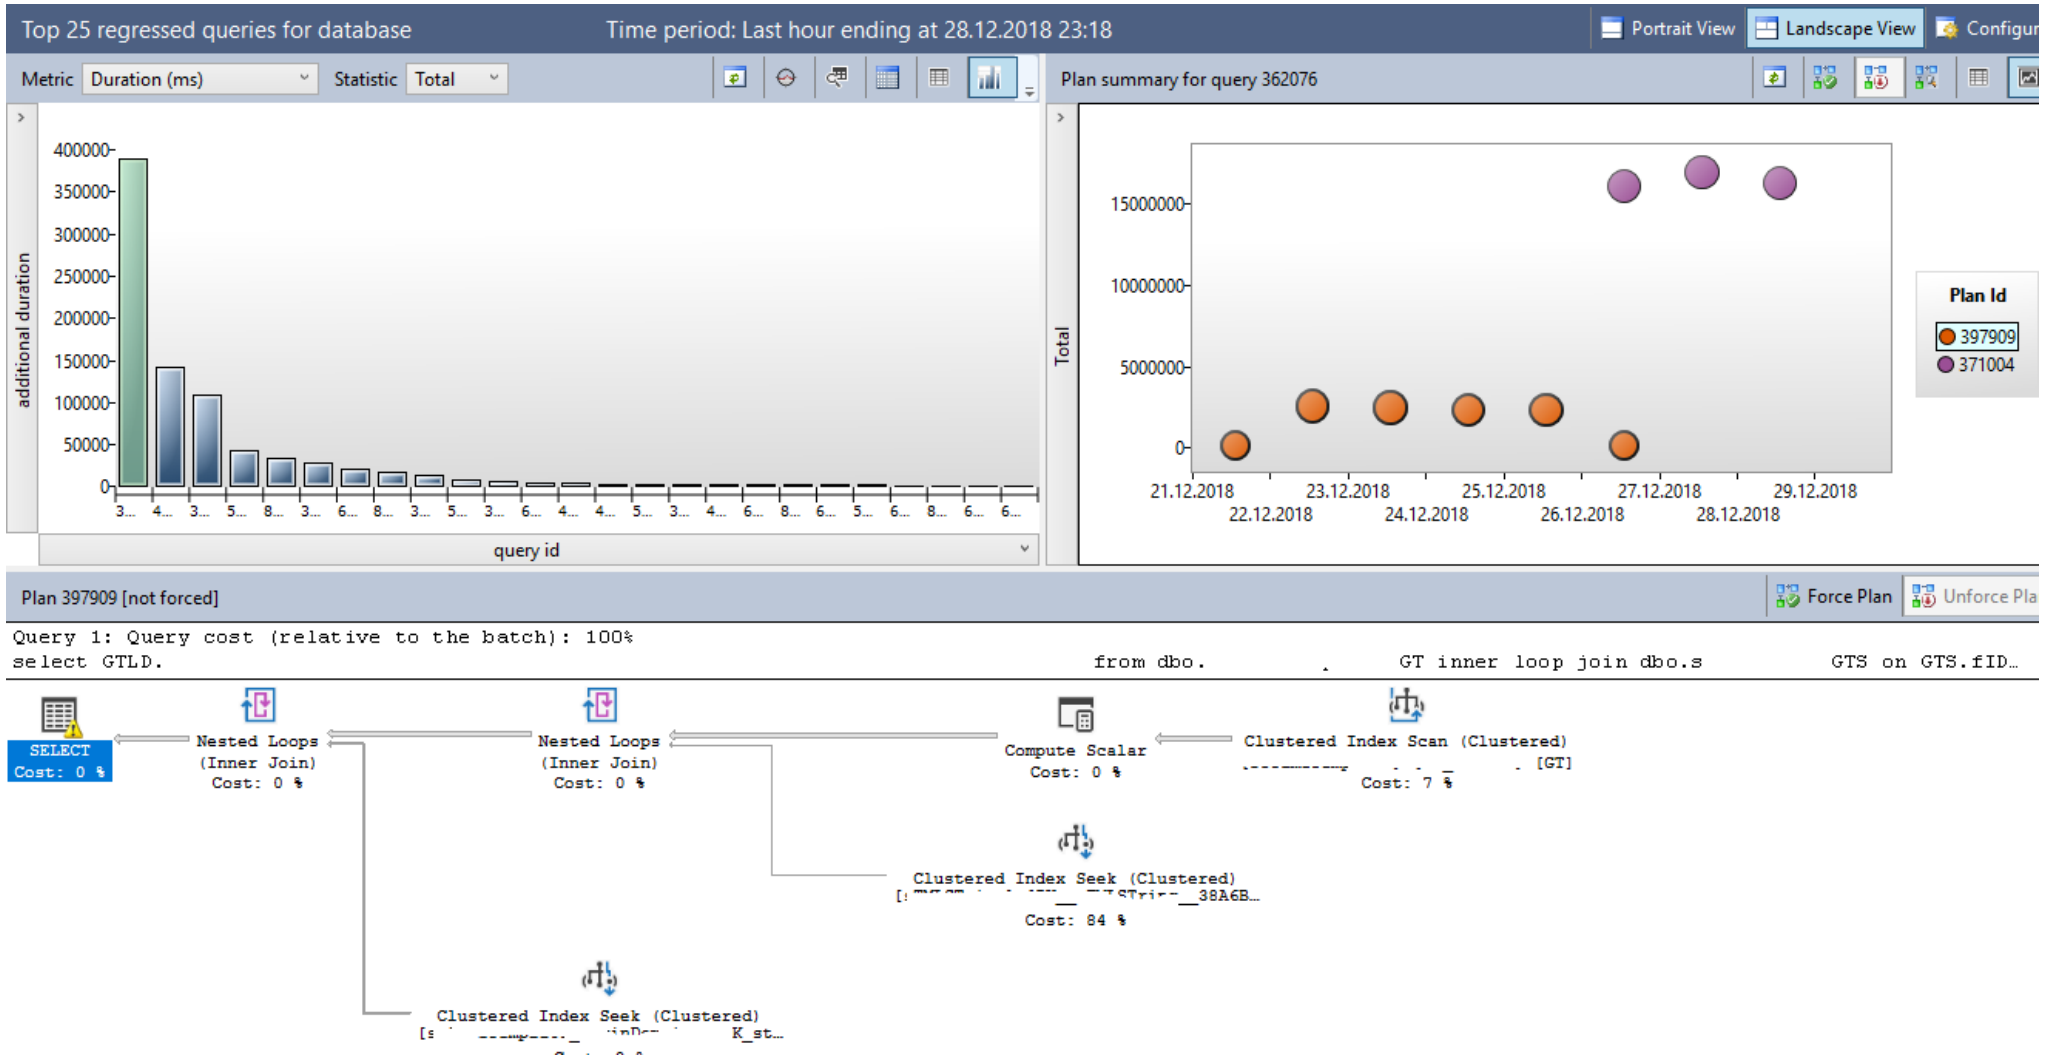

TOP RESOURCE CONSUMING QUERIES

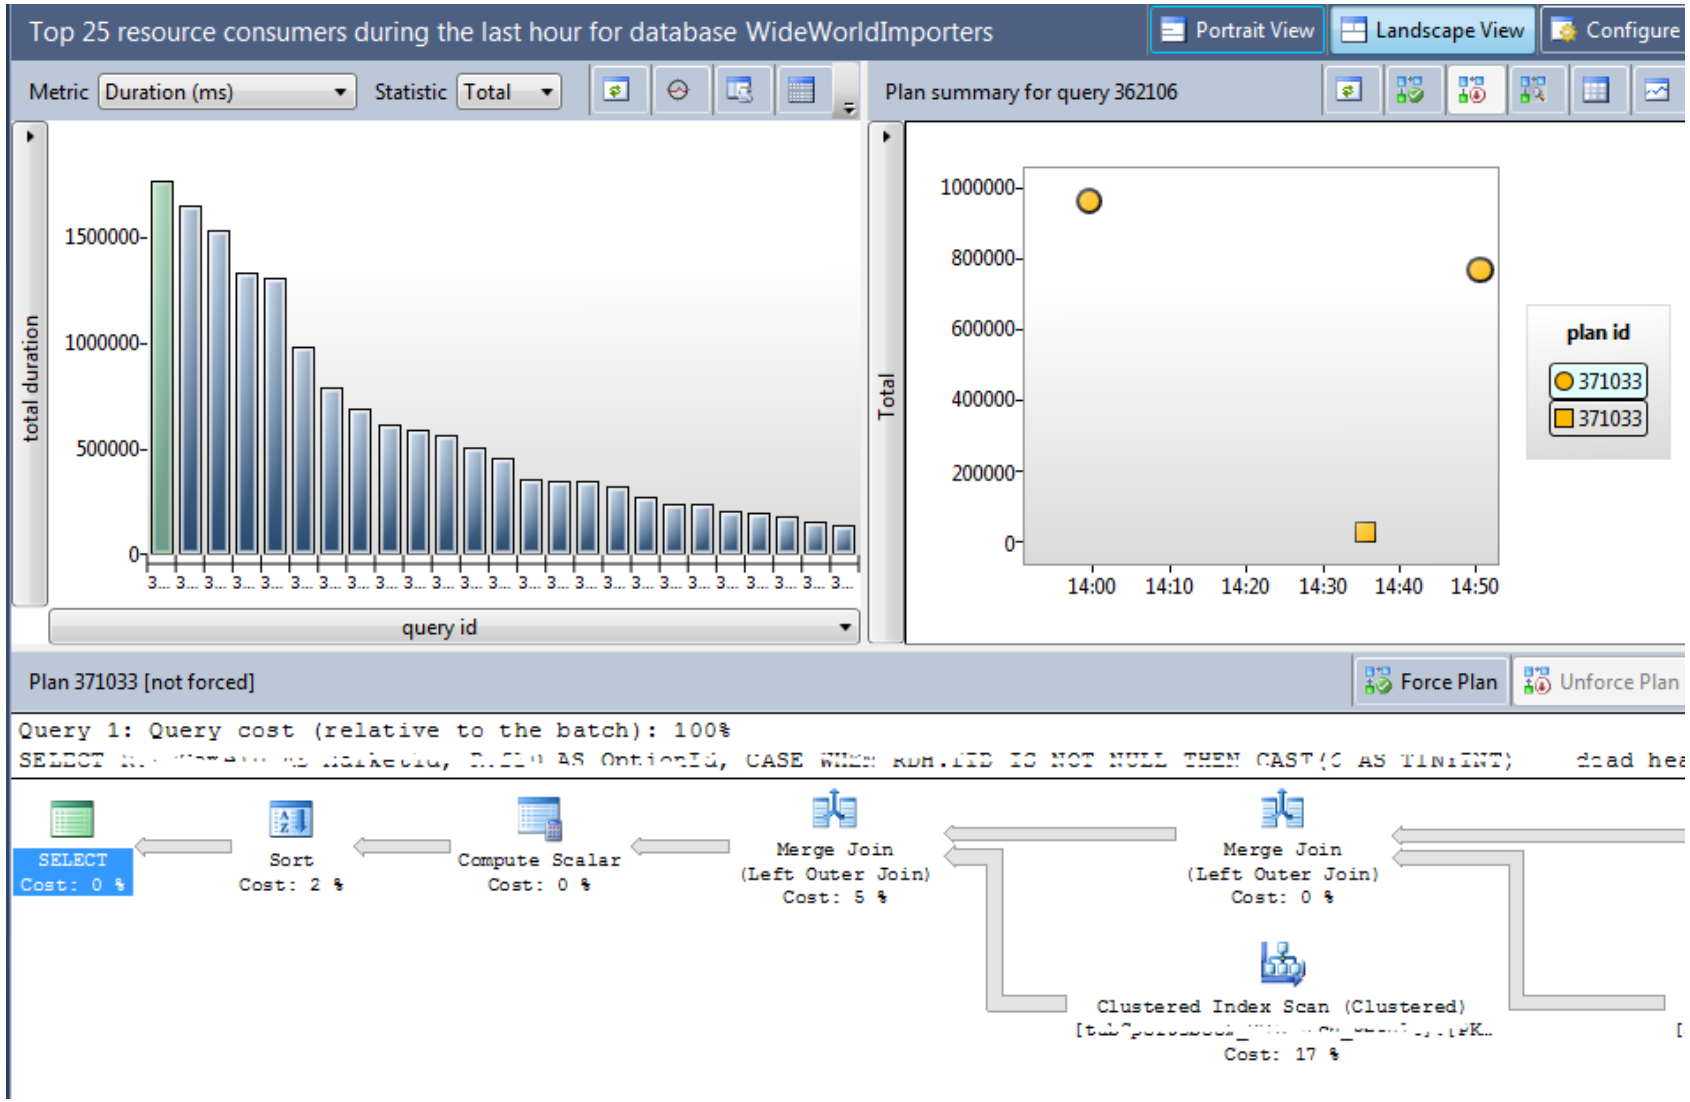

QUERIES WITH FORCED PLANS

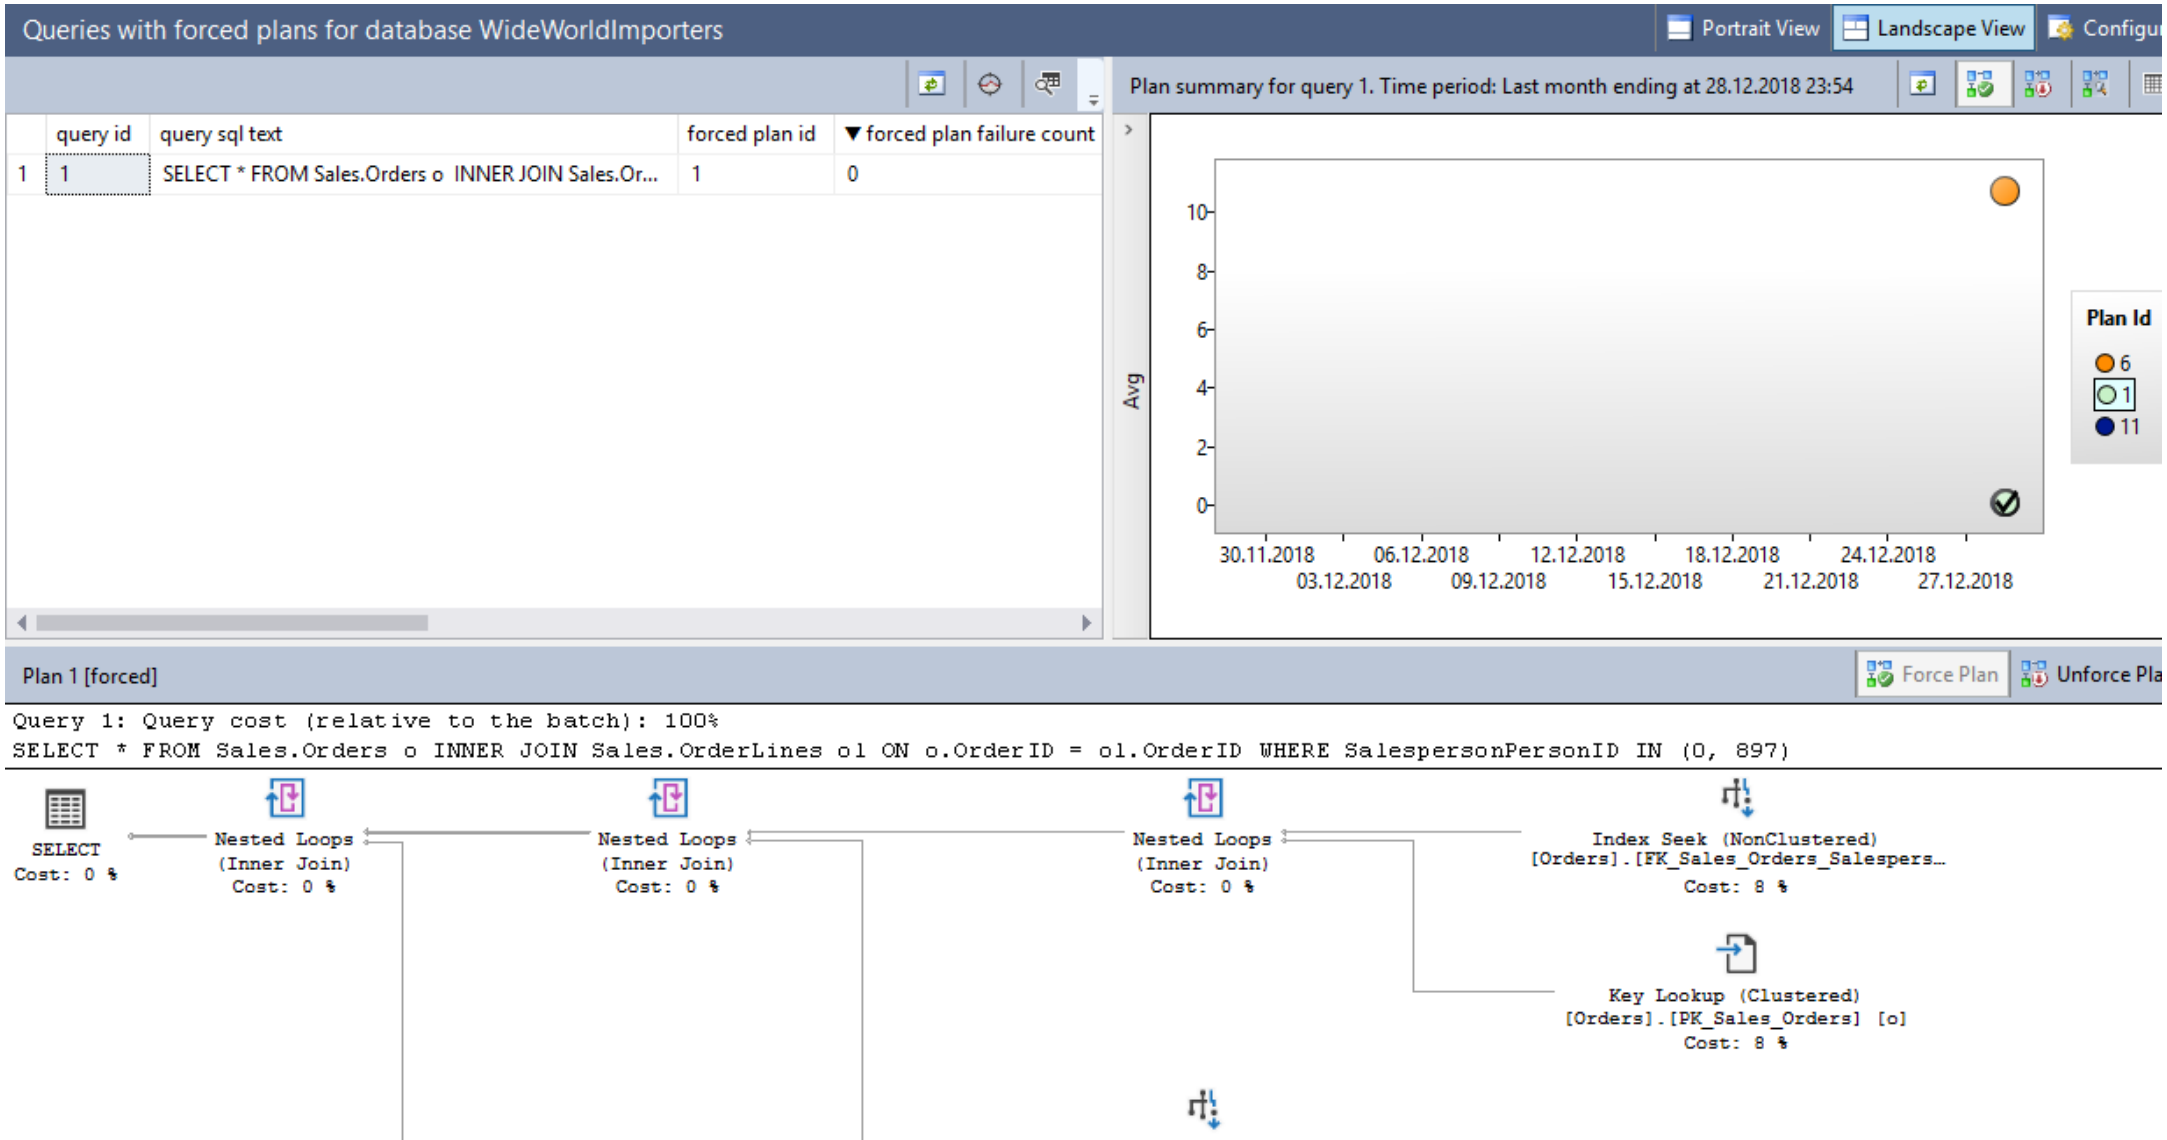

QUERY WAIT STATISTICS

Query wait statistics for database . Time period: Last week ending at 29.12.2018 00:23

Based on: total wait time

	wait category id	wait category	avg wait time	min wait time	max wait time	std dev wait time	▼ total wait time	execution count
1	3	Lock	1,49	0	4994	43,91	44185864	29721049

Query wait statistics for database . Time period: Last week ending at 29.12.2018 00:19

Portrait View | Landscape View | Configur

Based on: total wait time

avg wait time
max wait time
min wait time
std dev wait time
total wait time

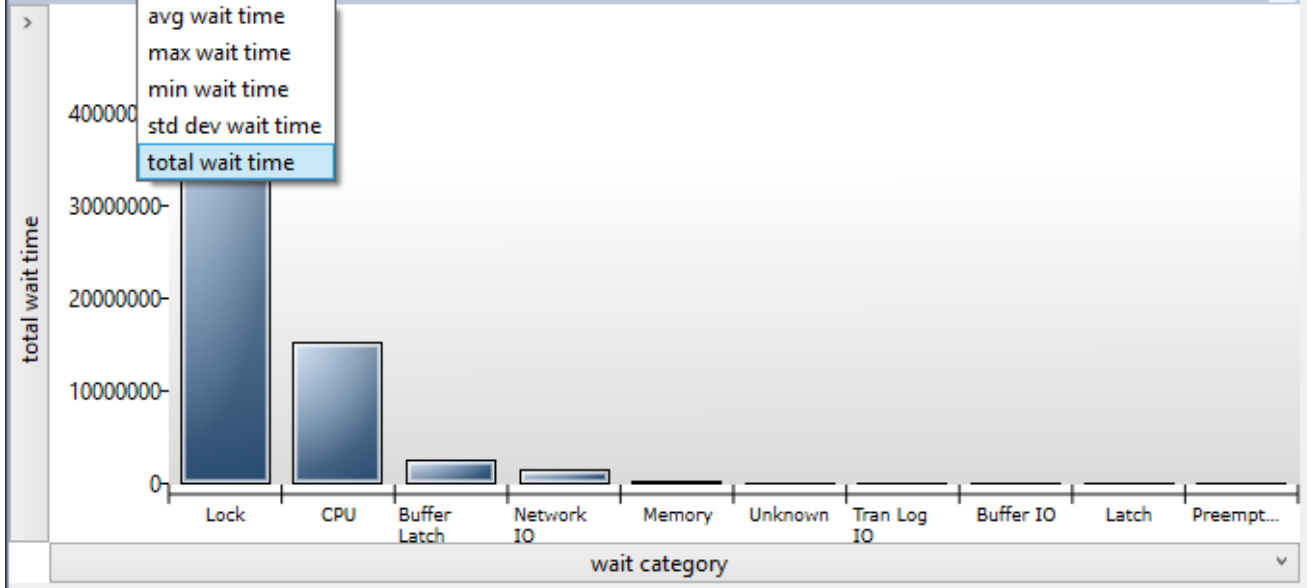

Plan

Force Plan | Unforce Pla

CAPTURING WAITS WITH QUERY STORE

CAPTURING WAITS

Wait category	Wait types
CPU	SOS_SCHEDULER_YIELD
Memory	RESOURCE_SEMAPHORE, CMEMTHREAD, CMEMPARTITIONED, EE_PMOLOCK, MEMORY_ALLOCATION_EXT, RESERVED_MEMORY_ALLOCATION_EXT, MEMORY_GRANT_UPDATE
Network IO	ASYNC_NETWORK_IO, NET_WAITFOR_PACKET, PROXY_NETWORK_IO, EXTERNAL_SCRIPT_NETWORK_IOF
Parallelism	CXPACKET, EXCHANGE
Lock	LCK_M_%
Latch	LATCH_%

QUERY STORE LIMITATIONS

- Option to store query store data in a filegroup other than PRIMARY
- DDL statements are not captured
- Tracking queries, where we changed the query (tuned it)
- GUI - Object view (SP)
- Enable Query Store for collection on a read-only replica in an Availability Group
- Exceptions are not caught for Hekaton

MS AZURE FEEDBACK ITEMS

278
votes

Voted!

Enable Query Store for collection on a read-only replica in an Availability Group

Currently, Query Store can only be enabled for the read-write database in Availability Group. As many customers issue queries against the read-only replicas, capturing query and performance metrics for those queries would be beneficial for understanding the workload, troubleshooting performance issues, etc.

41
votes

Voted!

Option to store query store data in a filegroup other than PRIMARY

Query store data is incredibly useful but depending on activity and settings it can grow to quite a large size. I would like an option to store this data on a filegroup other than PRIMARY. This would allow discretion regarding storage used, reduce the impact to recovery times and give DBAs more flexibility in managing query store data.

QUERY STORE ISSUES – LOAD AND UNLOAD

- **Loading**
  - By default, all queries are blocked from running in the database until Query Store is loaded
  - TF 7752 Query Store load asynchronously, and in the background will be in a read-only state until it is completely loaded, so queries can process while Query Store loads, but you will not be capturing them.
- **Shutdown**
  - shutdown is blocked by Query Store flushing process, you need to wait for recovery
  - TF 7745, you don't wait, but data loss
- **ALTER DATABASE … SET QUERY_STORE = OFF**
  - Here TF 7745 does not work!

QUERY STORE - RECOMMENDATIONS

- Control the size, do not let it grow
  - faster, less prone to load/shutdown issues
  - if you need more data, move it to your own tables
- Choose a runtime interval which is proper for your workload
- Do not use ALL for QUERY_CAPTURE_MODE
- Use AUTO for SIZE_BASED_CLEANUP_MODE
- Overhead (too many adhoc queries, too many queries, small interval too many plans, more impact on large QS
- Do not use DROP/CREATE and db rename or dbid()

QUERY STORE - RECOMMENDATIONS

- DBAs/DevOps
  - Identify the problem fast
  - Very fast workaround implementation (w/o BL knowledge, w/o code changes)
  - Learning about potential issues
- Developers
  - Learning about workload and execution plans
  - Analyze workload patterns
  - It‘s not a definitive solution, but it‘s a great help